

The PRISM Handbook

Seventeen documents covering how PRISM is built, deployed, operated and used. Every claim here was read out of the code it cites, and code excerpts are quoted rather than paraphrased so they stay checkable.

The documents

#	Document	What it answers
01	Architecture	What PRISM is, the three design commitments, the service catalogue, the request path, trust boundaries
02	Deployment architecture	Compose and Helm topologies, the edge, ports, volumes, the production posture, sizing
03	Module interaction	Who calls whom, how identity travels, timeout chains, what happens when a dependency is down
04	Running flows	Capture, conversion, stage change, import, reconciliation, intel, handover — traced end to end
05	Temporal workflows	The fourteen workflows, why a signal is untrusted, retry tiers, SLA and escalation
06	Code map	Directory by directory, plus a "where do I change X?" index
07	User management & RBAC	The authority model, twelve roles, two matrices, scoping, provisioning
08	The Register	The CRUD factory, concurrency, lifecycle enforcement, reconciliation, seeding
09	Backup & restore	What must survive, what is actually backed up, the restore drill, RPO/RTO
10	Upgrade & rollback	<code>prism-deploy.sh</code> in full, schema-change asymmetry, the failure playbook
11	ATLAS usage	The ten tabs as the desk meets them, what each role sees, common tasks
12	VocX & STT	The capture pipeline, timeout chain, recording cap, sizing, tuning
13	Operations	Health checks, logs, symptom → cause → fix, routine checks
14	Configuration	Every variable, compose vs Helm, the production checklist
15	Data model & ERD	Every table, the shared row spine, FK policy, RLS, migrations
16	The credit pipeline	The three gates, what approval and rejection each do, the walk-back, how lines and deals conclude
17	User management in practice	Role by role: what each can do and see, scoped vs full, maker-checker pairs, the employee use-cases, the live-roster caution
—	Publishing	Building the PDF and mirroring the handbook into Confluence

Reading paths

Security review



07 RBAC




03 Module interaction





01 Architecture §5







On call / deploying





New engineer





If you are...	Read, in order
A new backend engineer	01 → 06 → 08 → 15 → 03
A new front-end engineer	11 → 06 → 07 → 04
Deploying for the first time	02 → 14 → 10 → 09
On call	13 → 10 → 02

If you are...	Read, in order
Adding a business process	05 → 04 → 08
Training the desk	11 → 04
Doing a security review	07 → 03 → 01 §5 → 05 §4
Sizing the box	02 §7 → 12 §7

The five things everyone should know

Whatever your role, these five come up constantly.

1. **The Register is the only service that owns data.** Everything else is stateless and goes through its API. (01)

2. **Authorization is decided once at the gateway and carried in a signed token.** Nothing downstream trusts a plaintext identity header. (03)
3. **A Temporal signal is not authority.** The durable decision record in the Register is. (05)
4. **nginx caches upstream IPs.** Recreate a container by hand and you must reload the edge, or every page 502s while everything reports healthy. (13)
5. **Three paths never enter git or a delivery archive:** `deploy/compose/.env`, `deploy/vocx-secrets/`, `deploy/nginx/certs/`. (09)

Known gaps, stated plainly

These are real and worth knowing before you rely on something that is not there.

Gap	Where it is discussed
MinIO objects and the VocX audio archive have no scheduled backup. The database does.	09 §1, §7
RPO is up to 24 hours , and backups sit on the same host.	09 §8
The ATLAS UI has no automated test suite and no working eslint config. Verification is <code>tsc --noEmit</code> plus browser checks.	06 §10
<code>POST /v1/{subject}/{id}/interactions</code> is not idempotency-keyed — a retried write can duplicate a timeline row.	03 §7, 08 §7
A VocX-created lead is assigned to the capturer , not the company's BD owner.	04 §1
Dev ports (5432, 8000–8002, 9000/9001) are published by the same compose file used in production. Bind them to loopback or close them at the security group.	02 §1
LMS servicing is deferred (<code>LMS_ENABLED=false</code>). The book ends at <code>Disbursed</code> .	14 §2

Related documents outside the handbook

The handbook explains the system. These go deeper on specific subjects.

Document	Subject
docs/FOUNDATION_SPEC.md	The product specification
docs/SCHEMA.md	Column-level schema detail
docs/ARCHITECTURE.md	The earlier architecture note
docs/adr/	Architecture decision records (register-first, entity-centric, optimistic concurrency, monorepo, self-hosted Postgres, transient retry)
docs/STAGES_AND_APPROVALS.md	Stage and approval rules
docs/LENDING_WORKFLOW_DESIGN.md	The lending workflow in depth
docs/MIS_IMPORT.md	The spreadsheet import contract, sheet by sheet
docs/VOCX_API.md	The VocX API surface
docs/USER_PROVISIONING.md · docs/GOOGLE_SSO.md	Provisioning and SSO setup
docs/DEPLOYMENT.md · docs/PRODUCTION_DEPLOYMENT.md · docs/DEPLOY_RUNBOOK.md	Deployment guides
docs/E2E_RUNBOOK.md · docs/E2E_MASTER_FLOW.md · docs/UI_E2E_FLOW.md	End-to-end test journeys
docs/ATLAS_TAB_DB_API_MAP.md · docs/ATLAS_UI_FIELD_MAP.md	Tab → table → endpoint, and field-level maps
docs/openapi/	Exported OpenAPI specs
postman/	API collections and environments
BACKEND_STANDARDS.md · CONTRIBUTING.md	Coding conventions

Keeping this accurate

A handbook that drifts is worse than none, because people trust it. Two habits:

- **When you change something a document describes, update the document in the same commit.** Each document names the files it draws from; the "where do I change X?" table in [06](#) is the fastest way to find the affected page.
- **Quote code rather than paraphrasing it.** A quoted excerpt that has drifted is visibly wrong; a paraphrase just quietly becomes false.

Diagrams are Mermaid, so they render on GitHub and diff as text.